

Multivariate Statistical Analysis

Lecture 14

Fudan University

luoluo@fudan.edu.cn

- 1 Bayesian Multivariate Linear Regression
- 2 Principal Components Analysis
- 3 Principal Coordinate Analysis
- 4 Kernel Principal Component Analysis

- 1 Bayesian Multivariate Linear Regression
- 2 Principal Components Analysis
- 3 Principal Coordinate Analysis
- 4 Kernel Principal Component Analysis

Bayesian Multivariate Linear Regression

We can additionally suppose each b_{ij} independently follows

$$b_{ij} \sim \mathcal{N}(0, \tau^2),$$

then the posterior likelihood function is

$$\begin{aligned} & L(\mathbf{B}, \mathbf{\Sigma}) \\ &= \prod_{i=1}^N \frac{1}{\sqrt{(2\pi)^p \det(\mathbf{\Sigma})}} \exp \left(-\frac{1}{2} (\mathbf{B}^\top \mathbf{x}_i - \mathbf{y}_i)^\top \mathbf{\Sigma}^{-1} (\mathbf{B}^\top \mathbf{x}_i - \mathbf{y}_i) \right) \\ & \quad \cdot \prod_{i=1}^p \prod_{j=1}^q \frac{1}{\sqrt{2\pi\tau^2}} \exp \left(-\frac{b_{ij}^2}{2\tau^2} \right) \\ & \propto \frac{1}{(\det(\mathbf{\Sigma}))^{N/2}} \exp \left(-\frac{1}{2} \text{tr}((\mathbf{XB} - \mathbf{Y})\mathbf{\Sigma}^{-1}(\mathbf{XB} - \mathbf{Y})^\top) - \frac{1}{2\tau^2} \|\mathbf{B}\|_F^2 \right), \end{aligned}$$

which leads to

$$\text{vec}(\hat{\mathbf{B}}) = (\mathbf{I}_q \otimes \tau^2 \mathbf{X}^\top \mathbf{X} + \mathbf{\Sigma} \otimes \mathbf{I}_p)^{-1} \text{vec}(\tau^2 \mathbf{X}^\top \mathbf{Y}).$$

Bayesian Multivariate Linear Regression

We typically suppose

$$\beta_{(i)} \stackrel{\text{i.i.d}}{\sim} \mathcal{N}_q(\mathbf{0}, \tau^2 \mathbf{\Sigma}), \quad \text{where} \quad \mathbf{B} = \begin{bmatrix} \beta_{(1)}^\top \\ \vdots \\ \beta_{(p)}^\top \end{bmatrix} \in \mathbb{R}^{p \times q},$$

then the posterior likelihood function is

$$\begin{aligned} & L(\mathbf{B}, \mathbf{\Sigma}) \\ &= \prod_{i=1}^N \frac{1}{\sqrt{(2\pi)^p \det(\mathbf{\Sigma})}} \exp \left(-\frac{1}{2} (\mathbf{B}^\top \mathbf{x}_i - \mathbf{y}_i)^\top \mathbf{\Sigma}^{-1} (\mathbf{B}^\top \mathbf{x}_i - \mathbf{y}_i) \right) \\ & \quad \cdot \prod_{j=1}^p \frac{1}{\sqrt{(2\pi)^q \det(\mathbf{\Sigma})}} \exp \left(-\frac{1}{2\tau^2} \beta_{(j)}^\top \mathbf{\Sigma}^{-1} \beta_{(j)} \right) \\ & \propto \frac{1}{(\det(\mathbf{\Sigma}))^{N/2}} \exp \left(-\frac{1}{2} \text{tr} \left((\mathbf{XB} - \mathbf{Y}) \mathbf{\Sigma}^{-1} (\mathbf{XB} - \mathbf{Y})^\top \right) - \frac{1}{2\tau^2} \mathbf{B} \mathbf{\Sigma}^{-1} \mathbf{B}^\top \right). \end{aligned}$$

Bayesian Multivariate Linear Regression

We have

$$\hat{\mathbf{B}}_{\lambda} = (\mathbf{X}^{\top} \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^{\top} \mathbf{Y},$$

and

$$\hat{\boldsymbol{\Sigma}}_{\lambda} = \frac{1}{N} \mathbf{Y}^{\top} (\mathbf{I} - \mathbf{X}(\mathbf{X}^{\top} \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^{\top}) \mathbf{Y},$$

where $\lambda = 1/\tau^2$.

Outline

- 1 Bayesian Multivariate Linear Regression
- 2 Principal Components Analysis
- 3 Principal Coordinate Analysis
- 4 Kernel Principal Component Analysis

Population Principal Components Analysis

Let \mathbf{x} be a p -dimensional random vector with mean $\mathbf{0}$ and covariance matrix $\mathbf{\Sigma} \succ \mathbf{0}$.

Let $\mathbf{u}_1 \in \mathbb{R}^p$ with $\|\mathbf{u}_1\|_2 = 1$ and maximizing the variance of $\mathbf{u}_1^\top \mathbf{x}$, then

$$(\mathbf{\Sigma} - \lambda_1 \mathbf{I})\mathbf{u}_1 = \mathbf{0},$$

where λ_1 is the largest root of

$$\det(\mathbf{\Sigma} - \lambda \mathbf{I}) = 0.$$

- 1 We call $y_1 = \mathbf{u}_1^\top \mathbf{x}$ as the first principle component of \mathbf{x} .
- 2 The pair $\lambda_1 \in \mathbb{R}$ and $\mathbf{u}_1 \in \mathbb{R}^p$ are the largest eigenvalue and corresponding eigenvector of $\mathbf{\Sigma}$.

Population Principal Components Analysis

For the second principle components

$$y_2 = \mathbf{u}_2^\top \mathbf{x},$$

we determine $\mathbf{u}_2 \in \mathbb{R}^p$ by maximizing the variance of y_2 under the constraints $\|\mathbf{u}_2\|_2 = 1$ and y_2 be uncorrelated with y_1 .

For the k -th principle component

$$y_k = \mathbf{u}_k^\top \mathbf{x},$$

we determine \mathbf{u}_k by maximizing the variance of y_k under the constraints $\|\mathbf{u}_k\|_2 = 1$ and y_k be uncorrelated with y_1, \dots, y_{k-1} .

Population Principal Components Analysis

Let vector $\mathbf{u}_k \in \mathbb{R}^p$ the k -th principle component

$$y_k = \mathbf{u}_k^\top \mathbf{x}$$

holds that

$$(\mathbf{\Sigma} - \lambda_k \mathbf{I})\mathbf{u}_k = \mathbf{0},$$

where λ_k is the k -th largest root of

$$\det(\mathbf{\Sigma} - \lambda \mathbf{I}) = 0.$$

The pair $\lambda_k \in \mathbb{R}$ and $\mathbf{u}_k \in \mathbb{R}^p$ are the k -th largest eigenvalue and corresponding eigenvector of $\mathbf{\Sigma}$.

Population Principal Components Analysis

Let vector $\mathbf{u}_k \in \mathbb{R}^p$ the k -th principle component

$$y_k = \mathbf{u}_k^\top \mathbf{x}$$

holds that

$$(\mathbf{\Sigma} - \lambda_k \mathbf{I})\mathbf{u}_k = \mathbf{0},$$

where λ_k is the k -th largest root of

$$\det(\mathbf{\Sigma} - \lambda \mathbf{I}) = 0.$$

The pair $\lambda_k \in \mathbb{R}$ and $\mathbf{u}_k \in \mathbb{R}^p$ are the k -th largest eigenvalue and corresponding eigenvector of $\mathbf{\Sigma}$.

PCA for dimensionality Reduction

We can write

$$\mathbf{U}_k = [\mathbf{u}_1 \quad \cdots \quad \mathbf{u}_k] \in \mathbb{R}^{p \times k} \quad \text{and} \quad \mathbf{\Lambda}_k = \begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_k \end{bmatrix} \in \mathbb{R}^{k \times k}$$

contains the top- k eigenvectors and eigenvalues pairs of $\mathbf{\Sigma}$, that is

$$\mathbf{\Sigma} \mathbf{U}_k = \mathbf{U}_k \mathbf{\Lambda}_k \quad \text{with} \quad \mathbf{U}_k^\top \mathbf{U}_k = \mathbf{I}.$$

PCA for dimensionality Reduction

We can keep $\mathbf{U}_k \in \mathbb{R}^{p \times k}$ and transform $\mathbf{x} \in \mathbb{R}^p$ to

$$\mathbf{U}_k^\top \mathbf{x} \in \mathbb{R}^k,$$

where $k \ll p$.

The information of \mathbf{x} can be estimated by

$$\hat{\mathbf{x}} = \mathbf{U}_k (\mathbf{U}_k^\top \mathbf{x}) \in \mathbb{R}^p.$$

We have

$$\text{Cov}[\hat{\mathbf{x}}] = \mathbf{U}_k \mathbf{\Lambda}_k \mathbf{U}_k^\top,$$

which is the best rank- k approximation of $\mathbf{\Sigma}$.

Outline

- 1 Bayesian Multivariate Linear Regression
- 2 Principal Components Analysis
- 3 Principal Coordinate Analysis**
- 4 Kernel Principal Component Analysis

Sample Principal Components Analysis

Given observation $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^p$, we construct sample covariance

$$\mathbf{S} = \frac{1}{N-1} \sum_{\alpha=1}^N (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^\top, \quad \text{where } \bar{\mathbf{x}} = \frac{1}{N} \sum_{\alpha=1}^N \mathbf{x}_\alpha.$$

Let spectral decomposition of \mathbf{S} be $\mathbf{S} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}$, where $\mathbf{U} \in \mathbb{R}^{p \times p}$ is orthogonal and $\mathbf{\Lambda} \in \mathbb{R}^{p \times p}$ is diagonal.

We write

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{bmatrix} \in \mathbb{R}^{N \times p},$$

which results the sample principle components

$$\mathbf{Y} = \begin{bmatrix} (\mathbf{x}_1 - \bar{\mathbf{x}})^\top \mathbf{U}_k \\ \vdots \\ (\mathbf{x}_N - \bar{\mathbf{x}})^\top \mathbf{U}_k \end{bmatrix} = \mathbf{H}\mathbf{X}\mathbf{U}_k \in \mathbb{R}^{N \times k}, \quad \text{where } \mathbf{H} = \mathbf{I} - \frac{1}{N} \mathbf{1}_N \mathbf{1}_N^\top \in \mathbb{R}^{N \times N}.$$

Principal Coordinate Analysis

We consider the case of $p \geq N$ and define

$$\mathbf{T} = \frac{1}{N-1} \mathbf{H} \mathbf{X} \mathbf{X}^T \mathbf{H} \in \mathbb{R}^{N \times N}$$

with spectral decomposition

$$\mathbf{T} = \mathbf{V} \mathbf{\Gamma} \mathbf{V}^T,$$

where $\mathbf{V} \in \mathbb{R}^{N \times N}$ is orthogonal and $\mathbf{\Gamma} \in \mathbb{R}^{N \times N}$ is diagonal.

The matrix $\mathbf{Y} \in \mathbb{R}^{N \times k}$ can be written as

$$\mathbf{Y} = \mathbf{V}_k \mathbf{\Gamma}_k^{1/2} \in \mathbb{R}^{N \times k}.$$

Kernel Principal Component Analysis

We consider the case of $p \geq N$ and define

$$\mathbf{T} = \frac{1}{N-1} \mathbf{H} \mathbf{X} \mathbf{X}^T \mathbf{H} \in \mathbb{R}^{N \times N}$$

with spectral decomposition

$$\mathbf{T} = \mathbf{V} \mathbf{\Gamma} \mathbf{V}^T,$$

where $\mathbf{V} \in \mathbb{R}^{N \times N}$ is orthogonal and $\mathbf{\Gamma} \in \mathbb{R}^{N \times N}$ is diagonal.

The matrix $\mathbf{Y} \in \mathbb{R}^{N \times k}$ can be written as

$$\mathbf{Y} = \mathbf{V}_k \mathbf{\Gamma}_k^{1/2} \in \mathbb{R}^{N \times k}.$$

Outline

- 1 Bayesian Multivariate Linear Regression
- 2 Principal Components Analysis
- 3 Principal Coordinate Analysis
- 4 Kernel Principal Component Analysis

Kernel Principal Component Analysis

The matrix $\mathbf{S} \in \mathbb{R}^{p \times p}$ corresponds to outer product the centralized data, that is

$$\mathbf{X}^T \mathbf{H} \mathbf{H} \mathbf{X} = \sum_{\alpha=1}^N (\mathbf{x}_{\alpha} \mathbf{H})(\mathbf{x}_{\alpha} \mathbf{H})^T \in \mathbb{R}^{p \times p}.$$

The matrix $\mathbf{T} \in \mathbb{R}^{N \times N}$ corresponds to the inner product, that is

$$\mathbf{H} \mathbf{X} \mathbf{X}^T \mathbf{H} = \mathbf{H} \begin{bmatrix} \mathbf{x}_1^T \mathbf{x}_1 & \mathbf{x}_1^T \mathbf{x}_2 & \dots & \mathbf{x}_1^T \mathbf{x}_N \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{x}_N^T \mathbf{x}_1 & \mathbf{x}_N^T \mathbf{x}_2 & \dots & \mathbf{x}_N^T \mathbf{x}_N \end{bmatrix} \mathbf{H} \in \mathbb{R}^{N \times N}.$$

We achieve kernel PCA by replacing $\mathbf{X} \mathbf{X}^T$ with some kernel matrix.

Kernel Principal Component Analysis

We map the sample $\mathbf{x}_\alpha \in \mathcal{X} \subseteq \mathbb{R}^p$ to the feature space $\mathcal{H} \subseteq \mathbb{R}^d$, that is

$$\phi : \mathcal{X} \rightarrow \mathcal{H},$$

and define the kernel function (inner product)

$$K(\mathbf{x}, \mathbf{y}) \triangleq \phi(\mathbf{x})^\top \phi(\mathbf{y}).$$

We replace $\mathbf{X}\mathbf{X}^\top \in \mathbb{R}^{N \times N}$ with the kernel matrix

$$\begin{aligned} \mathbf{K} &= \begin{bmatrix} \phi(\mathbf{x}_1)^\top \\ \vdots \\ \phi(\mathbf{x}_N)^\top \end{bmatrix} [\phi(\mathbf{x}_1) \quad \cdots \quad \phi(\mathbf{x}_N)] \\ &= \begin{bmatrix} K(\mathbf{x}_1, \mathbf{x}_1) & K(\mathbf{x}_1, \mathbf{x}_2) & \cdots & K(\mathbf{x}_1, \mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ K(\mathbf{x}_N, \mathbf{x}_1) & K(\mathbf{x}_N, \mathbf{x}_2) & \cdots & K(\mathbf{x}_N, \mathbf{x}_N) \end{bmatrix} \in \mathbb{R}^{N \times N}, \end{aligned}$$

Kernel Principal Component Analysis

We replace

$$\mathbf{T} = \frac{1}{N-1} \mathbf{H} \mathbf{X} \mathbf{X}^{\top} \mathbf{H} \in \mathbb{R}^{N \times N}$$

with

$$\mathbf{T}_K = \frac{1}{N-1} \mathbf{H} \mathbf{K} \mathbf{H}$$

and achieve kernel PCA by spectral decomposition on \mathbf{T}_K .