

Multivariate Statistical Analysis

Lecture 15

Fudan University

luoluo@fudan.edu.cn

Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis
- 3 Kernel Principal Component Analysis
- 4 Factor Analysis
- 5 Probabilistic Principle Component Analysis
- 6 Course Summary

Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis
- 3 Kernel Principal Component Analysis
- 4 Factor Analysis
- 5 Probabilistic Principle Component Analysis
- 6 Course Summary

Principal Components Analysis

Let \mathbf{x} be a p -dimensional random vector with mean $\mathbf{0}$ and covariance matrix $\mathbf{\Sigma} \succ \mathbf{0}$.

Let $\mathbf{u}_1 \in \mathbb{R}^p$ with $\|\mathbf{u}_1\|_2 = 1$ and maximizing the variance of $\mathbf{u}_1^\top \mathbf{x}$, then

$$(\mathbf{\Sigma} - \lambda_1 \mathbf{I})\mathbf{u}_1 = \mathbf{0},$$

where λ_1 is the largest root of

$$\det(\mathbf{\Sigma} - \lambda \mathbf{I}) = 0.$$

- 1 We call $y_1 = \mathbf{u}_1^\top \mathbf{x}$ as the first principle component of \mathbf{x} .
- 2 The pair $\lambda_1 \in \mathbb{R}$ and $\mathbf{u}_1 \in \mathbb{R}^p$ are the largest eigenvalue and corresponding eigenvector of $\mathbf{\Sigma}$.

Principal Components Analysis

For the second principle components

$$y_2 = \mathbf{u}_2^\top \mathbf{x},$$

we determine $\mathbf{u}_2 \in \mathbb{R}^p$ by maximizing the variance of y_2 under the constraints $\|\mathbf{u}_2\|_2 = 1$ and y_2 be uncorrelated with y_1 .

For the k -th principle component

$$y_k = \mathbf{u}_k^\top \mathbf{x},$$

we determine \mathbf{u}_k by maximizing the variance of y_k under the constraints $\|\mathbf{u}_k\|_2 = 1$ and y_k be uncorrelated with y_1, \dots, y_{k-1} .

Principal Components Analysis

Let vector $\mathbf{u}_k \in \mathbb{R}^p$ the k -th principle component

$$y_k = \mathbf{u}_k^\top \mathbf{x}$$

holds that

$$(\mathbf{\Sigma} - \lambda_k \mathbf{I})\mathbf{u}_k = \mathbf{0},$$

where λ_k is the k -th largest root of

$$\det(\mathbf{\Sigma} - \lambda \mathbf{I}) = 0.$$

The pair $\lambda_k \in \mathbb{R}$ and $\mathbf{u}_k \in \mathbb{R}^p$ are the k -th largest eigenvalue and corresponding eigenvector of $\mathbf{\Sigma}$.

Principal Components Analysis

Let vector $\mathbf{u}_k \in \mathbb{R}^p$ the k -th principle component

$$y_k = \mathbf{u}_k^\top \mathbf{x}$$

holds that

$$(\mathbf{\Sigma} - \lambda_k \mathbf{I})\mathbf{u}_k = \mathbf{0},$$

where λ_k is the k -th largest root of

$$\det(\mathbf{\Sigma} - \lambda \mathbf{I}) = 0.$$

The pair $\lambda_k \in \mathbb{R}$ and $\mathbf{u}_k \in \mathbb{R}^p$ are the k -th largest eigenvalue and corresponding eigenvector of $\mathbf{\Sigma}$.

PCA for dimensionality Reduction

We can write

$$\mathbf{U}_k = [\mathbf{u}_1 \quad \cdots \quad \mathbf{u}_k] \in \mathbb{R}^{p \times k} \quad \text{and} \quad \mathbf{\Lambda}_k = \begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_k \end{bmatrix} \in \mathbb{R}^{k \times k}$$

contains the top- k eigenvectors and eigenvalues pairs of $\mathbf{\Sigma}$, that is

$$\mathbf{\Sigma} \mathbf{U}_k = \mathbf{U}_k \mathbf{\Lambda}_k \quad \text{with} \quad \mathbf{U}_k^\top \mathbf{U}_k = \mathbf{I}.$$

PCA for dimensionality Reduction

We can keep $\mathbf{U}_k \in \mathbb{R}^{p \times k}$ and transform $\mathbf{x} \in \mathbb{R}^p$ to

$$\mathbf{U}_k^\top \mathbf{x} \in \mathbb{R}^k,$$

where $k \ll p$.

The information of \mathbf{x} can be estimated by

$$\hat{\mathbf{x}} = \mathbf{U}_k (\mathbf{U}_k^\top \mathbf{x}) \in \mathbb{R}^p.$$

We have

$$\text{Cov}[\hat{\mathbf{x}}] = \mathbf{U}_k \mathbf{\Lambda}_k \mathbf{U}_k^\top,$$

which is the best rank- k approximation of $\mathbf{\Sigma}$.

Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis**
- 3 Kernel Principal Component Analysis
- 4 Factor Analysis
- 5 Probabilistic Principle Component Analysis
- 6 Course Summary

Sample Principal Components Analysis

Given observation $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^p$, we construct sample covariance

$$\mathbf{S} = \frac{1}{N-1} \sum_{\alpha=1}^N (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^\top, \quad \text{where } \bar{\mathbf{x}} = \frac{1}{N} \sum_{\alpha=1}^N \mathbf{x}_\alpha.$$

Let spectral decomposition of \mathbf{S} be $\mathbf{S} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}$, where $\mathbf{U} \in \mathbb{R}^{p \times p}$ is orthogonal and $\mathbf{\Lambda} \in \mathbb{R}^{p \times p}$ is diagonal.

We write

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{bmatrix} \in \mathbb{R}^{N \times p},$$

which results the sample principle components

$$\mathbf{Y} = \begin{bmatrix} (\mathbf{x}_1 - \bar{\mathbf{x}})^\top \mathbf{U}_k \\ \vdots \\ (\mathbf{x}_N - \bar{\mathbf{x}})^\top \mathbf{U}_k \end{bmatrix} = \mathbf{H}\mathbf{X}\mathbf{U}_k \in \mathbb{R}^{N \times k}, \quad \text{where } \mathbf{H} = \mathbf{I} - \frac{1}{N} \mathbf{1}_N \mathbf{1}_N^\top \in \mathbb{R}^{N \times N}.$$

Principal Coordinate Analysis

We consider the case of $p \geq N$ and define

$$\mathbf{T} = \frac{1}{N-1} \mathbf{H} \mathbf{X} \mathbf{X}^T \mathbf{H} \in \mathbb{R}^{N \times N}$$

with spectral decomposition

$$\mathbf{T} = \mathbf{V} \mathbf{\Gamma} \mathbf{V}^T,$$

where $\mathbf{V} \in \mathbb{R}^{N \times N}$ is orthogonal and $\mathbf{\Gamma} \in \mathbb{R}^{N \times N}$ is diagonal.

The matrix $\mathbf{Y} \in \mathbb{R}^{N \times k}$ can be written as

$$\mathbf{Y} = \mathbf{V}_k \mathbf{\Gamma}_k^{1/2} \in \mathbb{R}^{N \times k}.$$

Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis
- 3 Kernel Principal Component Analysis**
- 4 Factor Analysis
- 5 Probabilistic Principle Component Analysis
- 6 Course Summary

Kernel Principal Component Analysis

We map the sample $\mathbf{x}_\alpha \in \mathcal{X} \subseteq \mathbb{R}^p$ to the feature space $\mathcal{H} \subseteq \mathbb{R}^d$, that is

$$\phi : \mathcal{X} \rightarrow \mathcal{H},$$

and define the corresponding kernel function (inner product)

$$K(\mathbf{x}, \mathbf{y}) \triangleq \phi(\mathbf{x})^\top \phi(\mathbf{y}).$$

Kernel Principal Component Analysis

The matrix

$$\mathbf{T} = \frac{1}{N-1} \mathbf{H} \mathbf{X} \mathbf{X}^\top \mathbf{H} \in \mathbb{R}^{N \times N}$$

contains

$$\mathbf{H} \mathbf{X} \mathbf{X}^\top \mathbf{H} = \mathbf{H} \begin{bmatrix} \mathbf{x}_1^\top \mathbf{x}_1 & \mathbf{x}_1^\top \mathbf{x}_2 & \dots & \mathbf{x}_1^\top \mathbf{x}_N \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{x}_N^\top \mathbf{x}_1 & \mathbf{x}_N^\top \mathbf{x}_2 & \dots & \mathbf{x}_N^\top \mathbf{x}_N \end{bmatrix} \mathbf{H} \in \mathbb{R}^{N \times N}.$$

We replace the inner product $\mathbf{x}_i^\top \mathbf{x}_j$ with

$$K(\mathbf{x}_i, \mathbf{x}_j) \triangleq \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j).$$

Kernel Principal Component Analysis

We replace $\mathbf{X}\mathbf{X}^\top \in \mathbb{R}^{N \times N}$ with the kernel matrix

$$\mathbf{K} = \begin{bmatrix} K(\mathbf{x}_1, \mathbf{x}_1) & K(\mathbf{x}_1, \mathbf{x}_2) & \dots & K(\mathbf{x}_1, \mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ K(\mathbf{x}_N, \mathbf{x}_1) & K(\mathbf{x}_N, \mathbf{x}_2) & \dots & K(\mathbf{x}_N, \mathbf{x}_N) \end{bmatrix} \in \mathbb{R}^{N \times N}$$

and replace $\mathbf{T} \in \mathbb{R}^{N \times N}$ with

$$\mathbf{T}_K = \frac{1}{N-1} \mathbf{H} \mathbf{K} \mathbf{H}.$$

The kernel PCA is achieved by spectral decomposition on \mathbf{T}_K .

Kernel Principal Component Analysis

We replace $\mathbf{X}\mathbf{X}^\top \in \mathbb{R}^{N \times N}$ with the kernel matrix

$$\mathbf{K} = \begin{bmatrix} K(\mathbf{x}_1, \mathbf{x}_1) & K(\mathbf{x}_1, \mathbf{x}_2) & \dots & K(\mathbf{x}_1, \mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ K(\mathbf{x}_N, \mathbf{x}_1) & K(\mathbf{x}_N, \mathbf{x}_2) & \dots & K(\mathbf{x}_N, \mathbf{x}_N) \end{bmatrix} \in \mathbb{R}^{N \times N}$$

and replace $\mathbf{T} \in \mathbb{R}^{N \times N}$ with

$$\mathbf{T}_K = \frac{1}{N-1} \mathbf{H} \mathbf{K} \mathbf{H}.$$

The kernel PCA is achieved by spectral decomposition on \mathbf{T}_K .

Kernel Principal Component Analysis

Examples of kernel functions:

- 1 We define the polynomial kernel as

$$K(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^\top \mathbf{y} + c)^d$$

for some $c \in \mathbb{R}$ and $d \in \mathbb{N}$.

- 2 We define the Gaussian kernel (radial basis function kernel) as

$$K(\mathbf{x}, \mathbf{y}) = \exp \left(-\frac{\|\mathbf{x} - \mathbf{y}\|_2^2}{2\sigma^2} \right).$$

Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis
- 3 Kernel Principal Component Analysis
- 4 Factor Analysis**
- 5 Probabilistic Principle Component Analysis
- 6 Course Summary

Factor Analysis

Let the observable vector $\mathbf{y} \in \mathbb{R}^p$ be written as

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \boldsymbol{\mu} + \boldsymbol{\epsilon},$$

where

- ① $\mathbf{W} \in \mathbb{R}^{p \times q}$ is the loading matrix (parameter),
- ② $\mathbf{x} \in \mathbb{R}^q$ is the common factor (parameter/random vector),
- ③ $\boldsymbol{\mu} \in \mathbb{R}^p$ is the mean vector (parameter),
- ④ $\boldsymbol{\epsilon} \in \mathbb{R}^p$ is the specific factor (random vector).

The model is similar to regression, but \mathbf{x} is unobserved.

Factor Analysis

Example of sports games:

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mu + \epsilon.$$

- 1 \mathbf{y} : performance in real-world
- 2 \mathbf{W} : system of the game
- 3 \mathbf{x} : attributes in the game
- 4 μ : average attributes
- 5 ϵ : noise/exception

Yao Ming's Stats			
57 Games	82 Games	80 Games	48 Games
22.3 Points	17.5 Points	18.3 Points	25.0 Points
51.9% FG%	52.2% FG%	55.2% FG%	51.6% FG%
0.0% 3P%	0.0% 3P%	0.0% 3P%	0.0% 3P%
85.3% FT%	80.9% FT%	86.2% FT%	86.2% FT%
10.2 Rebounds	9.4 Rebounds	9.4 Rebounds	9.4 Rebounds
1.5 Assists	2.0 Assists	2.0 Assists	2.0 Assists
0.5 Steals	0.4 Steals	0.4 Steals	0.4 Steals
1.6 Blocks	2.0 Blocks	2.0 Blocks	2.0 Blocks
2.6 Turnovers	2.5 Turnovers	3.5 Turnovers	3.5 Turnovers
25.6 PER	21.9 PER	26.5 PER	26.5 PER
0.211 ws/48	0.202 ws/48	0.220 ws/48	0.220 ws/48



Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis
- 3 Kernel Principal Component Analysis
- 4 Factor Analysis
- 5 Probabilistic Principle Component Analysis**
- 6 Course Summary

Probabilistic Principle Component Analysis

Let $\mathbf{y}_1, \dots, \mathbf{y}_N \in \mathbb{R}^p$ be N independent observations and we have

$$\mathbf{y}_\alpha = \mathbf{W}\mathbf{x}_\alpha + \boldsymbol{\mu} + \epsilon_\alpha,$$

where

$$\mathbf{x}_\alpha \sim \mathcal{N}_q(\mathbf{0}, \mathbf{I}) \quad \text{and} \quad \epsilon_\alpha \sim \mathcal{N}_p(\mathbf{0}, \sigma^2 \mathbf{I})$$

are independent for some $\sigma^2 > 0$ and $q < \min\{N, p\}$.

We target to estimate parameters

$$\mathbf{W} \in \mathbb{R}^{p \times q}, \quad \boldsymbol{\mu} \in \mathbb{R}^p \quad \text{and} \quad \sigma \in (0, +\infty)$$

by maximum likelihood estimation for given $\mathbf{y}_1, \dots, \mathbf{y}_N$.

Probabilistic Principle Component Analysis

Consider that

$$\mathbf{y}_\alpha \sim \mathcal{N}_p(\boldsymbol{\mu}, \mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}).$$

We construct the likelihood function

$$\begin{aligned} & L(\boldsymbol{\mu}, \mathbf{W}, \sigma^2) \\ &= \prod_{\alpha=1}^N \frac{1}{\sqrt{(2\pi)^p \det(\boldsymbol{\Sigma})}} \exp\left(-\frac{1}{2}(\mathbf{y}_\alpha - \boldsymbol{\mu})^\top (\mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I})^{-1}(\mathbf{y}_\alpha - \boldsymbol{\mu})\right), \end{aligned}$$

then we have

$$\begin{aligned} & \ln L(\boldsymbol{\mu}, \mathbf{W}, \sigma^2) \\ & \propto -\frac{N}{2} \ln \det(\mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I}) - \frac{1}{2} \sum_{\alpha=1}^N (\mathbf{y}_\alpha - \boldsymbol{\mu})^\top (\mathbf{W}\mathbf{W}^\top + \sigma^2\mathbf{I})^{-1}(\mathbf{y}_\alpha - \boldsymbol{\mu}). \end{aligned}$$

The Maximum Likelihood Estimators

The maximum likelihood estimators of μ , \mathbf{W} and σ^2 are

$$\hat{\mu} = \bar{\mathbf{y}} = \frac{1}{N} \sum_{\alpha=1}^N \mathbf{y}_{\alpha}, \quad \hat{\mathbf{W}} = \mathbf{U}_q (\mathbf{\Lambda}_q - \hat{\sigma}^2 \mathbf{I}) \mathbf{R} \quad \text{and} \quad \hat{\sigma}^2 = \frac{1}{p-q} \sum_{j=q+1}^p \lambda_j,$$

where

- 1 $\mathbf{\Lambda}_q \in \mathbb{R}^{q \times q}$ is diagonal with the largest q eigenvalues $\lambda_1, \dots, \lambda_q$ of

$$\hat{\mathbf{\Sigma}} = \frac{1}{N} \sum_{\alpha=1}^N (\mathbf{y}_{\alpha} - \bar{\mathbf{y}})(\mathbf{y}_{\alpha} - \bar{\mathbf{y}})^{\top};$$

- 2 $\mathbf{U}_q \in \mathbb{R}^{p \times q}$ is orthogonal column consisting of the eigenvectors associate with $\lambda_1, \dots, \lambda_q$;
- 3 $\mathbf{R} \in \mathbb{R}^{q \times q}$ is any orthogonal matrix.

The Maximum Likelihood Estimators

The maximum likelihood estimators also minimize the error with respect to Frobenius norm

$$(\hat{\mathbf{W}}, \hat{\sigma}^2) = \arg \min_{\mathbf{W} \in \mathbb{R}^{p \times q}, \sigma^2 \in \mathbb{R}^+} \left\| \hat{\mathbf{\Sigma}} - (\mathbf{W}\mathbf{W}^\top + \sigma^2 \mathbf{I}) \right\|_F.$$

The Expectation-Maximization Algorithm

For the model

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \boldsymbol{\mu} + \epsilon,$$

where $\mathbf{x} \sim \mathcal{N}_q(\mathbf{0}, \mathbf{I})$ and $\epsilon \sim \mathcal{N}_p(\mathbf{0}, \sigma^2 \mathbf{I})$ are independent.

We regard $\{\mathbf{x}_\alpha\}_{\alpha=1}^N$ as missing data and $\{\mathbf{x}_\alpha, \mathbf{y}_\alpha\}_{\alpha=1}^N$ as the complete data, then we can achieve

$$\mathbf{y}_\alpha \mid \mathbf{x}_\alpha \sim \mathcal{N}_p(\mathbf{W}\mathbf{x}_\alpha + \boldsymbol{\mu}, \sigma^2 \mathbf{I})$$

and

$$\mathbf{x}_\alpha \mid \mathbf{y}_\alpha \sim \mathcal{N}_q(\mathbf{M}^{-1} \mathbf{W}^\top (\mathbf{y}_\alpha - \boldsymbol{\mu}), \sigma^2 \mathbf{M}^{-1}),$$

where $\mathbf{M} = \mathbf{W}^\top \mathbf{W} + \sigma^2 \mathbf{I}$.

The Expectation-Maximization Algorithm

The update of the EM algorithm

- 1 In E-step, we take the expectation

$$l_C = \mathbb{E} \left[\ln \left(\prod_{\alpha=1}^N f(\mathbf{x}_\alpha | \mathbf{y}_\alpha) \right) \right].$$

- 2 In the M-step, we maximized l_C with respect to \mathbf{W} and σ^2 :

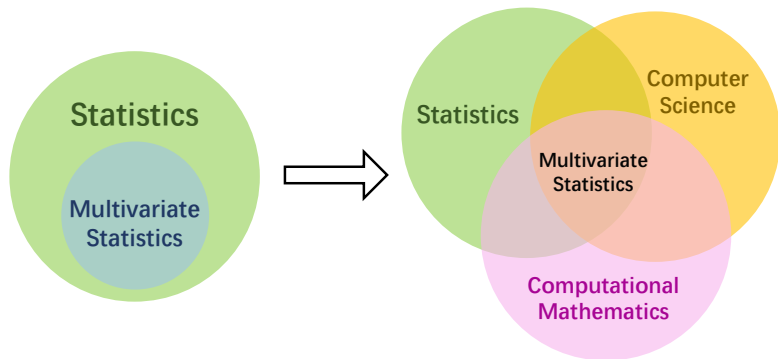
$$\begin{aligned} \mathbf{W}_+ &= \hat{\Sigma} \mathbf{W} (\sigma^2 \mathbf{I} + \mathbf{M}^{-1} \mathbf{W}^\top \hat{\Sigma} \mathbf{W})^{-1}, \\ \sigma_+^2 &= \frac{1}{p} \text{tr} \left(\hat{\Sigma} - \hat{\Sigma} \mathbf{W} \mathbf{M}^{-1} \mathbf{W}_+^\top \right). \end{aligned}$$

Note that the computational complexity of EM is $\mathcal{O}(Npq)$, while the spectral decomposition in MLE requires $\mathcal{O}(Np^2 + p^3)$.

Outline

- 1 Principal Components Analysis
- 2 Principal Coordinate Analysis
- 3 Kernel Principal Component Analysis
- 4 Factor Analysis
- 5 Probabilistic Principle Component Analysis
- 6 Course Summary**

Multivariate Statistics



Good Luck on Finals!

Multivariate Statistical Analysis (DATA 13004)

Final Exam



Matrix Calculus



Linear Algebra



Machine Learning



Statistics



Optimization

