

Optimization Theory

Lecture 04

Fudan University

luoluo@fudan.edu.cn

- 1 Second-Order Characterization
- 2 Black Box Model
- 3 Gradient Descent Methods

1 Second-Order Characterization

2 Black Box Model

3 Gradient Descent Methods

Second-Order Characterization

Theorem (Smoothness and Convexity)

Let $f(\cdot)$ be a twice differentiable function defined on \mathbb{R}^d

- ① It is L -smooth if and only if $-L\mathbf{I} \preceq \nabla^2 f(\mathbf{x}) \preceq L\mathbf{I}$ for all $\mathbf{x} \in \mathbb{R}^d$.
- ② It is convex if and only if $\nabla^2 f(\mathbf{x}) \succeq \mathbf{0}$ for all $\mathbf{x} \in \mathbb{R}^d$.
- ③ It is μ -strongly-convex if and only if $\nabla^2 f(\mathbf{x}) \succeq \mu\mathbf{I}$ for all $\mathbf{x} \in \mathbb{R}^d$.

Sometimes, we say $f(\cdot)$ is ℓ -weakly convex if the function

$$g(\mathbf{x}) = f(\mathbf{x}) + \frac{\ell}{2} \|\mathbf{x}\|_2^2$$

is convex for some $\ell > 0$.

Theorem

Let $f : \mathbb{R}^d \rightarrow \mathbb{R}$ be a twice differentiable function. Suppose that $\nabla^2 f(\cdot)$ is continuous in an open neighborhood of $\mathbf{x}^* \in \mathbb{R}^d$.

① If \mathbf{x}^* is a local minimizer of $f(\cdot)$, then it holds that

$$\nabla f(\mathbf{x}^*) = \mathbf{0} \quad \text{and} \quad \nabla^2 f(\mathbf{x}^*) \succeq \mathbf{0}.$$

② If it holds that

$$\nabla f(\mathbf{x}^*) = \mathbf{0} \quad \text{and} \quad \nabla^2 f(\mathbf{x}^*) \succ \mathbf{0},$$

then the point \mathbf{x}^* is a strict local minimizer of $f(\cdot)$.

Second-Order Characterization

Some examples:

- ① For unconstrained quadratic problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) \triangleq \frac{1}{2} \mathbf{x}^\top \mathbf{A} \mathbf{x} - \mathbf{b}^\top \mathbf{x},$$

where $\mathbf{A} \in \mathbb{R}^{d \times d}$. We have

$$\nabla^2 f(\mathbf{x}) = \mathbf{A}.$$

- ② For regularized generalized linear model

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) \triangleq \frac{1}{n} \sum_{i=1}^n \phi_i(\mathbf{a}_i^\top \mathbf{x}) + \frac{\lambda}{2} \|\mathbf{x}\|_2^2.$$

where $\phi_i : \mathbb{R}^d \rightarrow \mathbb{R}$ is twice differentiable. We have

$$\nabla f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \phi'_i(\mathbf{a}_i^\top \mathbf{x}) \mathbf{a}_i + \lambda \mathbf{x} \quad \text{and} \quad \nabla^2 f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \phi''_i(\mathbf{a}_i^\top \mathbf{x}) \mathbf{a}_i \mathbf{a}_i^\top + \lambda \mathbf{I}.$$

1 Second-Order Characterization

2 Black Box Model

3 Gradient Descent Methods

Convergence Criteria

For the unconstrained convex optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}),$$

the convergence of an algorithm can be measured by the following in metrics:

- 1 Convergence in parameter (suppose there exists optimal solution \mathbf{x}^*), where we measure the distance

$$\|\mathbf{x}_t - \mathbf{x}^*\|_2.$$

- 2 Convergence of objective value, measured by objective suboptimality

$$f(\mathbf{x}_t) - \inf_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}).$$

- 3 Convergence of gradient

$$\|\nabla f(\mathbf{x}_t)\|_2.$$

Convergence Criteria

If $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is smooth and convex and has an optimal solution \mathbf{x}^* , then

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \leq \langle \nabla f(\mathbf{x}^*), \mathbf{x}_t - \mathbf{x}^* \rangle + \frac{L}{2} \|\mathbf{x}_t - \mathbf{x}^*\|_2^2 = \frac{L}{2} \|\mathbf{x}_t - \mathbf{x}^*\|_2^2,$$

and

$$\|\nabla f(\mathbf{x}_t)\|_2 = \|\nabla f(\mathbf{x}_t) - \nabla f(\mathbf{x}^*)\|_2 \leq L \|\mathbf{x}_t - \mathbf{x}^*\|_2,$$

which implies convergence in parameter implies convergence in objective value and gradient.

The reverse directions may not hold if the objective is not strongly-convex.

Black Box Model

Local black box:

- 1 The only information available for the numerical scheme is the answer of the oracle.
- 2 The oracle is local.

Different types of oracle:

- 1 Zero-order oracle: returns the function value $f(\mathbf{x})$.
- 2 First-order oracle: returns the function value $f(\mathbf{x})$ and the gradient $\nabla f(\mathbf{x})$.
- 3 Second-order oracle: returns $f(\mathbf{x})$, $\nabla f(\mathbf{x})$, and the Hessian $\nabla^2 f(\mathbf{x})$.

Black Box Model

There are two participants in the black box model: a learner and an oracle.

- ① The learner has
 - infinite computational power,
 - knowledge of the function class to which f belongs,
 - knowledge of the domain.
- ② The oracle has specific knowledge of the function.

The key question:

How many queries to the oracles are necessary and sufficient to find an ϵ -approximate solution?

We will study this question from two perspectives:

- ① Upper bound: Designing algorithms.
- ② Lower bound: Information theoretic reasoning.

The strength of the black-box model:

- ① It will allow us to derive a complete theory of convex optimization.
- ② We will obtain matching upper and lower bounds on the oracle complexity for various subclasses of interesting functions.

The weakness of the black-box model:

- ① It does not limit our computational resources.
- ② The side information of the algorithm is ignored.

Outline

1 Second-Order Characterization

2 Black Box Model

3 Gradient Descent Methods

Gradient Descent Methods

We consider the optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}),$$

where $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is convex and L -smooth.

The gradient descent method

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta_t \nabla f(\mathbf{x}_t)$$

with $\eta_t = \eta \leq 1/L$ leads to

$$\frac{1}{T} \sum_{t=1}^T f(\mathbf{x}_t) \leq f(\hat{\mathbf{x}}) + \frac{L \|\mathbf{x}_0 - \hat{\mathbf{x}}\|_2^2}{2T}$$

for any $\hat{\mathbf{x}} \in \mathbb{R}^d$.

Minimizing Convex Function

The gradient descent method

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta \nabla f(\mathbf{x}_t)$$

with $\eta_t = \eta \leq 1/L$ leads to

$$\frac{1}{T} \sum_{t=1}^T f(\mathbf{x}_t) \leq f(\hat{\mathbf{x}}) + \frac{L \|\mathbf{x}_0 - \hat{\mathbf{x}}\|_2^2}{2T}$$

for any $\hat{\mathbf{x}} \in \mathbb{R}^d$.

Suppose $f(\cdot)$ has a minimizer \mathbf{x}^* and let $\bar{\mathbf{x}} = \frac{1}{T} \sum_{t=0}^{T-1} \mathbf{x}_t$, then we need

$$T \geq \left\lceil \frac{L \|\mathbf{x}_0 - \mathbf{x}^*\|_2^2}{2} \cdot \frac{1}{\epsilon} \right\rceil$$

to guarantee $f(\bar{\mathbf{x}}) - f(\mathbf{x}^*) \leq \epsilon$.

Last-Iterate Convergence

It is also possible to establish the last-iterate convergence

$$f(\mathbf{x}_T) - f(\mathbf{x}^*) \leq \frac{2L \|\mathbf{x}_0 - \mathbf{x}^*\|_2^2}{t + 4},$$

which is sublinear.

The proof depends on the results

$$\frac{1}{L} \|\nabla f(\mathbf{x}) - \nabla f(\mathbf{x}^*)\|_2^2 \leq \langle \nabla f(\mathbf{x}) - \nabla f(\mathbf{x}^*), \mathbf{x} - \mathbf{x}^* \rangle.$$

and

$$f(\mathbf{x}_{t+1}) \leq f(\mathbf{x}_t) - \frac{\eta}{2} \|\nabla f(\mathbf{x}_t)\|_2^2.$$

Nonconvex Optimization

The following inequality

$$f(\mathbf{x}_{t+1}) \leq f(\mathbf{x}_t) - \frac{\eta}{2} \|\nabla f(\mathbf{x}_t)\|_2^2.$$

does not depend on the convexity.

We uniformly sample $\hat{\mathbf{x}}$ from $\{\mathbf{x}_0, \dots, \mathbf{x}_{T-1}\}$, then

$$\mathbb{E} \|\nabla f(\hat{\mathbf{x}})\|_2^2 \leq \frac{2L(f(\mathbf{x}_0) - f^*)}{T},$$

where we suppose $f^* = \inf_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) > -\infty$

We require

$$T \geq \left\lceil \frac{2L(f(\mathbf{x}_0) - f^*)}{\epsilon^2} \right\rceil$$

to find an ϵ -stationary point of f in expectation.

Minimizing Strongly Convex Function

We consider using gradient descent method

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta \nabla f(\mathbf{x}_t)$$

with $\eta_t = \eta \leq 1/L$ to solve the optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}),$$

where $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is μ -strongly-convex and L -smooth.

It holds linear convergence rate

$$f(\mathbf{x}_T) - f(\mathbf{x}^*) \leq \left(1 - \frac{\mu}{L}\right)^T (f(\mathbf{x}_0) - f(\mathbf{x}^*)).$$

We require

$$T \geq \left\lceil \kappa \ln \left(\frac{f(\mathbf{x}_T) - f(\mathbf{x}^*)}{\epsilon} \right) \right\rceil$$

to guarantee $f(\mathbf{x}_T) - f(\mathbf{x}^*) \leq \epsilon$, where $\kappa \triangleq L/\mu$ is the condition number.

Example: Regularized Generalized Linear Model

For regularized linear regression

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) \triangleq \frac{1}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2 + \frac{\beta}{2} \|\mathbf{x}\|_2^2$$

where $\mathbf{A} \in \mathbb{R}^{n \times d}$, $\mathbf{b} \in \mathbb{R}^d$ and $\lambda > 0$.

We have

$$\nabla^2 f(\mathbf{x}) = \mathbf{A}^\top \mathbf{A} + \beta \mathbf{I} \quad \text{and} \quad \kappa = \frac{\lambda_1(\mathbf{A}^\top \mathbf{A}) + \beta}{\lambda_d(\mathbf{A}^\top \mathbf{A}) + \beta}.$$

Example: Quadratic Problem

We consider using gradient descent method to solve quadratic problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) \triangleq \frac{1}{2} \mathbf{x}^\top \mathbf{A} \mathbf{x} - \mathbf{b}^\top \mathbf{x},$$

where \mathbf{A} is positive definite.

Then we have

$$f(\mathbf{x}_T) - f(\mathbf{x}^*) \leq \left(1 - \frac{\lambda_d(\mathbf{A})}{\lambda_1(\mathbf{A})}\right)^T (f(\mathbf{x}_0) - f(\mathbf{x}^*)),$$

where $\lambda_1(\mathbf{A})$ and $\lambda_d(\mathbf{A})$ are the largest and the smallest eigenvalues of \mathbf{A} .

For positive semi-definite \mathbf{A} , what about the convergence rate?