Optimization Theory

Lecture 08

Fudan University

luoluo@fudan.edu.cn

Outline

1. Subgradient Descent Method
2. Smoothing Technique
3. Proximal Gradient Methods

Outline

1 Subgradient Descent Method

2 Smoothing Technique

3 Proximal Gradient Methods

Subgradient Descent Method

We consider optimization with a nonsmooth objective function

$$\min_{\mathbf{x}\in\mathcal{C}} f(\mathbf{x}).$$

Here we assume that $f:\mathbb{R}^d\to\mathbb{R}$ is G-Lipschitz and convex defined on a convex and closed set $\mathcal{C}\subseteq\mathbb{R}^d$, but not necessarily smooth.

We have introduced the subgradient method

$$\begin{cases}\tilde{\mathbf{x}}_{t+1} = \mathbf{x}_t - \eta_t \mathbf{g}_t,\\ \mathbf{x}_{t+1} = \mathrm{proj}_{\mathcal{C}}\left(\tilde{\mathbf{x}}_{t+1}\right).\end{cases}$$

Let $R = \sup_{\mathbf{x}\in\mathcal{C}} \|\mathbf{x}-\mathbf{x}_0\|_2$.

1. For convex case, it requires $\mathcal{O}(G^2R^2\epsilon^{-2})$ iterations.
2. For μ-strongly-convex, it requires $\mathcal{O}(G^2\mu^{-1}\epsilon^{-1})$ iterations.

Optimality of Subgradient Descent Method

Theorem

Given $G > 0$, $\mu > 0$, $d > t \geq 1$ and $\epsilon > 0$, there exists a G-Lipschitz and μ-strongly convex function $f : \mathbb{R}^d \to \mathbb{R}$ on

$$\mathcal{C} = \left\{ \mathbf{x} \in \mathbb{R}^d : \|\mathbf{x}\|_2 \leq \frac{G}{2\mu} \right\},$$

such that a first order optimization algorithm with initial point $\mathbf{x}_0 = 0$ can only produce solutions that satisfy

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \geq \frac{G^2}{8\mu(t+1)} - \epsilon,$$

where $\mathbf{x}^*$ is the solution of $\min_{\mathbf{x} \in \mathcal{C}} f(\mathbf{x})$.

Optimality of Subgradient Descent Method

Let $\mathbf{x} = [x_1, \dots, x_d] \in \mathbb{R}^d$ and define

$$f(\mathbf{x}) = \frac{G}{2} \max_{i=1,2\dots,t+1} \left(x_i + \frac{\epsilon}{i} \right) + \frac{\mu}{2} \|\mathbf{x}\|_2^2 .$$

1. The function f is G-Lipschitz continuous.
2. Any subgradient $\mathbf{g} \in \partial f(\mathbf{x})$ satisfies $\mathbf{g} = \lambda \mathbf{x} + 0.5G\mathbf{y}$, where

$$\mathbf{y} \in \mathrm{conv} \left\{ \mathbf{e}_i : x_i = \max_{k=1,2,\dots,t+1} x_j \right\} .$$

   We use $\mathrm{conv}(\mathcal{S})$ to present the convex hull of $\mathcal{S}$, which is defined as

$$\mathrm{conv}(\mathcal{S}) = \left\{ \sum_{i=1}^{m} \alpha_i \mathbf{x}_i : \mathbf{x}_i \in \mathcal{S}, \alpha_i \geq 0, \sum_{i=1}^{m} \alpha_i = 1 \right\} .$$

3. Check the zero-chain property.

Optimality of Subgradient Descent Method

For convex case, we can show the optimality by considering

$$f(\mathbf{x}) = \frac{G}{2} \max_{i=1,2\ldots,t+1} \left(x_i + \frac{\epsilon}{i} \right) + \frac{\mu}{2} \|\mathbf{x}\|_2^2 .$$

with

$$\mu = \frac{G}{2R\sqrt{t+1}}.$$

Outline

1. Subgradient Descent Method
2. Smoothing Technique
3. Proximal Gradient Methods

Smoothing Technique

Definition

We say the function $\tilde{f} : \mathbb{R}^d \to \mathbb{R}$ is an (L, ϵ)-smooth approximation of function $f : \mathbb{R}^d \to \mathbb{R}$ if $\tilde{f}$ is L-smooth and we have

$$\tilde{f}(\mathbf{x}) \leq f(\mathbf{x}) \leq \tilde{f}(\mathbf{x}) + \epsilon.$$

for all $\mathbf{x} \in \mathbb{R}^d$.

We can find approximate solution $\tilde{\mathbf{x}}$ for $\min_{\mathbf{x}\in\mathbb{R}^d} \tilde{f}(\mathbf{x})$ such that

$$\tilde{f}(\tilde{\mathbf{x}}) \leq \inf_{\mathbf{x}\in\mathbb{R}^d} \tilde{f}(\mathbf{x}) + \tilde{\epsilon},$$

then

$$\begin{aligned} f(\tilde{\mathbf{x}}) \leq& \tilde{f}(\tilde{\mathbf{x}}) + \epsilon \leq \inf_{\mathbf{x}\in\mathbb{R}^d} \tilde{f}(\mathbf{x}) + \tilde{\epsilon} + \epsilon \\ \leq& \inf_{\mathbf{x}\in\mathbb{R}^d} f(\mathbf{x}) + \tilde{\epsilon} + \epsilon. \end{aligned}$$

Smoothing Technique

Theorem

If $f : \mathbb{R}^d \to \mathbb{R}$ is convex and G-Lipschitz continuous, then

$$\tilde{f}(\mathbf{x}) = \min_{\mathbf{z}\in\mathbb{R}^d} \left(f(\mathbf{z}) + \frac{L}{2}\|\mathbf{z}-\mathbf{x}\|_2^2 \right).$$

is convex and it is a $(L, G^2/(2L))$-smooth approximation of $f(\mathbf{x})$.

Applying AGD to minimize $\tilde{f}(\mathbf{x})$ with $L = \mathcal{O}(G^2/\epsilon)$ can find $\mathcal{O}(\epsilon)$ suboptimal solution of $\min_{\mathbf{x}\in\mathbb{R}^d} f(\mathbf{x})$.

1. For convex $f(\mathbf{x})$, we require $\tilde{\mathcal{O}}(G/\epsilon)$ iterations.
2. For μ-strongly convex $f(\mathbf{x})$, we require $\tilde{\mathcal{O}}(G/\sqrt{\mu\epsilon})$ iterations.

Outline

1. Subgradient Descent Method
2. Smoothing Technique
3. Proximal Gradient Methods

Composite Convex Optimization Problem

We consider the problem of the form

$$\min_{\mathbf{x}\in\mathbb{R}^d} \phi(\mathbf{x}) \triangleq f(\mathbf{x}) + g(\mathbf{x}),$$

where $f : \mathbb{R}^d \to \mathbb{R}$ is a smooth convex function and $g : \mathbb{R}^d \to \mathbb{R}$ is convex but possibly nonsmooth.

1. We focus on the case that g has some simple form.
2. Subgradient method leads to slow convergence.
3. How to obtain the convergence rate like (accelerated) gradient descent?

Proximal Operator

We introduce the proximal operator as follows

$$\mathrm{prox}_h(\mathbf{x}) = \operatorname*{arg\,min}_{\mathbf{z}\in\mathbb{R}^d} \left(\frac{1}{2}\|\mathbf{z}-\mathbf{x}\|_2^2 + h(\mathbf{x})\right),$$

where $h : \mathbb{R}^d \to \mathbb{R}$ is convex but possible nonsmooth.

Proximal Operator

Recall that optimizing smooth convex function $f(\mathbf{x})$ by gradient descent

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \frac{1}{L}\nabla f(\mathbf{x}_t)$$

is based on minimizing RHS of

$$f(\mathbf{y}) \leq f(\mathbf{x}_t) + \langle \nabla f(\mathbf{x}_t), \mathbf{y} - \mathbf{x}_t \rangle + \frac{L}{2} \|\mathbf{y} - \mathbf{x}_t\|_2^2 .$$

Proximal Gradient Descent

For composite problem

$$\min_{\mathbf{x}\in\mathbb{R}^d} \phi(\mathbf{x}) \triangleq f(\mathbf{x}) + g(\mathbf{x}),$$

we can minimize RHS of

$$\phi(\mathbf{y}) = f(\mathbf{y}) + g(\mathbf{y}) \leq f(\mathbf{x}_t) + \langle \nabla f(\mathbf{x}_t), \mathbf{y} - \mathbf{x}_t \rangle + \frac{L}{2}\|\mathbf{y} - \mathbf{x}_t\|_2^2 + g(\mathbf{y}).$$

That is

$$\mathbf{x}_{t+1} = \mathrm{prox}_{\eta g}(\mathbf{x}_t - \eta \nabla f(\mathbf{x}_t)) \qquad \text{with} \qquad \eta = 1/L.$$

Proximal Gradient Descent

It can be computed efficiently for some simple $g(\cdot)$. For example:

1. Let $g(\mathbf{x}) = \lambda \|\mathbf{x}\|_1$, then

$$\mathrm{prox}_{\eta g}(\mathbf{x}) = \begin{bmatrix} \mathrm{sign}(x_1) \max\{|x_1| - \eta\lambda, 0\} \\ \mathrm{sign}(x_2) \max\{|x_2| - \eta\lambda, 0\} \\ \vdots \\ \mathrm{sign}(x_d) \max\{|x_d| - \eta\lambda, 0\} \end{bmatrix},$$

   which can be computed efficiently.

2. Let $g(\mathbf{x}) = \mathbb{1}_{\mathcal{C}}(\mathbf{x})$ for some closed convex $\mathcal{C}$, then

$$\mathrm{prox}_{\eta g}(\mathbf{x}) = \mathrm{proj}_{\mathcal{C}}(\mathbf{x}),$$

   which leads to

$$\mathbf{x}_{t+1} = \mathrm{proj}_{\mathcal{C}}(\mathbf{x}_t - \eta \nabla f(\mathbf{x}_t)).$$

Gradient Mapping

For function $\phi = f + g$ with convex functions $f : \mathbb{R}^d \to \mathbb{R}$, $g : \mathbb{R}^d \to \mathbb{R}$ and $\eta > 0$, we define the gradient mapping as follows

$$\mathcal{G}_{\eta g, f}(\mathbf{x}) = \frac{1}{\eta}(\mathbf{x} - \mathrm{prox}_{\eta g}(\mathbf{x} - \eta \nabla f(\mathbf{x}))),$$

which is a generalization of gradient operator $\nabla f(\mathbf{x})$.

The proximal gradient method

$$\mathbf{x}_{t+1} = \mathrm{prox}_{\eta g}(\mathbf{x}_t - \eta \nabla f(\mathbf{x}_t))$$

is equivalent to

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta \mathcal{G}_{\eta g, f}(\mathbf{x}_t).$$

Gradient Mapping

We consider the composite convex problem

$$\min_{\mathbf{x}\in\mathbb{R}^d} \phi(\mathbf{x}) \triangleq f(\mathbf{x}) + g(\mathbf{x}),$$

where $f : \mathbb{R}^d \to \mathbb{R}$ is L-smooth and convex and $g : \mathbb{R}^d \to \mathbb{R}$ is convex but possibly nonsmooth.

Let $\mathbf{x}^+ = \mathrm{prox}_{\eta g}(\mathbf{x} - \eta\nabla f(\mathbf{x}))$.

1. The point $\mathbf{x}^*$ is an optimal solution if and only if $\mathcal{G}_{\eta g,f}(\mathbf{x}^*) = \mathbf{0}$.
2. Suppose g is μ_g-strongly convex and $\eta < 2/(L-\mu)$, then

$$\|\mathcal{G}_{\eta g,f}(\mathbf{x})\|_2^2 \leq \frac{2/\eta}{2-\eta(L-\mu_g)}(\phi(\mathbf{x}) - \phi(\mathbf{x}^+)).$$

3. Suppose ϕ is μ_ϕ-strongly convex and $\eta < 1/L$, then

$$\phi(\mathbf{x}^+) \leq \phi(\mathbf{x}^*) + \frac{1}{2\mu_\phi}\|\mathcal{G}_{\eta g,f}(\mathbf{x})\|_2^2.$$

Convergence Analysis (Convex)

We consider the composite convex problem

$$\min_{\mathbf{x}\in\mathbb{R}^d} \phi(\mathbf{x}) \triangleq f(\mathbf{x}) + g(\mathbf{x}),$$

where $f:\mathbb{R}^d \to \mathbb{R}$ is L-smooth and convex and $g:\mathbb{R}^d \to \mathbb{R}$ is convex.

The proximal gradient method with $\eta = 1/L$ holds that

$$\phi(\mathbf{x}_T) \leq \phi(\mathbf{x}^*) + \frac{L}{2T}\left\|\mathbf{x}_0 - \mathbf{x}^*\right\|_2^2.$$

Additionally suppose ϕ is μ_ϕ-strongly convex leads to

$$\phi(\mathbf{x}_T) - \phi(\mathbf{x}^*) \leq \left(1 - \frac{\mu_\phi}{L + \mu_\phi}\right)^T (\phi(\mathbf{x}_0) - \phi(\mathbf{x}^*)).$$

Convergence Analysis (Nonconvex)

If we only suppose g is convex but allow f be nonconvex, then

$$\mathbb{E}\,\|\mathcal{G}_{\eta g,f}(\hat{\mathbf{x}})\|_2^2 \le \frac{2L(\phi(\mathbf{x}_0)-\phi^*)}{T},$$

where $\phi^* = \inf_{\mathbf{x}\in\mathbb{R}^d}\phi(\mathbf{x}) > -\infty$ and $\hat{\mathbf{x}}$ is uniformly sampled from

$$\{\mathbf{x}_0,\dots,\mathbf{x}_{T-1}\}.$$

Here, we say $\hat{\mathbf{x}}$ is an ϵ-stationary point of $\phi(\mathbf{x}) = f(\mathbf{x}) + g(\mathbf{x})$ if

$$\|\mathcal{G}_{\eta g,f}(\hat{\mathbf{x}})\|_2 \le \epsilon.$$

Accelerated Proximal Gradient Descent

We can also apply Nesterov's acceleration to proximal gradient methods

$$\mathbf{y}_t = \mathbf{x}_t + \beta_t(\mathbf{x}_t - \mathbf{x}_{t-1}),$$
$$\mathbf{x}_{t+1} = \mathrm{prox}_{\eta g}(\mathbf{y}_t - \eta_t \nabla f(\mathbf{y}_t)).$$

For convex case, it holds

$$\phi(\mathbf{x}_T) - \phi(\mathbf{x}^*) \leq \mathcal{O}\left(\frac{L}{T^2}\right).$$

For strongly-convex case, it holds

$$\phi(\mathbf{x}_T) - \phi(\mathbf{x}^*) \leq \mathcal{O}\left(\left(1 - \sqrt{\frac{\mu_\phi}{L}}\right)^T\right).$$

Subgradient Method vs. Proximal Gradient Method

Solving the composite convex problem

$$\min_{\mathbf{x}\in\mathbb{R}^d} \phi(\mathbf{x}) \triangleq f(\mathbf{x}) + g(\mathbf{x}),$$

by subgradient method are based on

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta_t(\nabla f(\mathbf{x}_t) + \boldsymbol{\xi}_t),$$

where $\boldsymbol{\xi}_t \in \partial g(\mathbf{x}_t)$.

The proximal gradient method is more progressive, since it holds that

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta_t(\nabla f(\mathbf{x}_t) + \boldsymbol{\xi}_{t+1}),$$

where $\boldsymbol{\xi}_{t+1} \in \partial g(\mathbf{x}_{t+1})$.